

# Homework Two

Robert George Phillips (rogphill@ucsc.edu)


February 21, 2018

**Question 1 (2 points):** Heaps naturally lead to a sorting algorithm, Heapsort. Starting from array  $A$ , build it into a min-heap. Repeatedly called Extract-Min to get the elements of  $A$  in sorted order. Give pseudocode for Heapsort and show that it is not stable.

**Analysis:** We start with an array  $A = \{A_0, A_1, \dots, A_k\}$  with  $k$  elements. We can turn this into a heap by iterating through each element and forming a tree structure where each node has two children. We can see this more easily with a sample array:

40	$30_1$	$30_2$	70	50	10	20
----	--------	--------	----	----	----	----

This sample array can simply be visualized into a heap, such that:



We have visualized a heap structure, but need to turn it into a min-heap, where the minimum element of the heap, 10, is at the root node. We can use min-heapify for this.

Min-heapify assumes that the subtree rooted at  $i$  has children that are heaps, but may not be a heap itself, and uses the fact that our base case is a single node that has no children – which is why we start at  $n - 1$  and decrement to 0. Additionally, if the element  $A[i].key$  is at least the minimum of the children’s keys, we are done. Otherwise, min-heapify will find a child  $j$  of  $A[i]$  with a minimum key and swap  $A[i]$  with  $A[j]$ . Finally, it will call itself recursively. In this way, the minimum value floats up and the larger values float down. In pseudocode:

```
Min-Heapify(A, i): //running time:  $O(\log n)$ 
    if A[i] has no children, return


    if A[i].key <= A[i].leftchild.key && A[i].rightchild.key, return

    else, find child j of A[i] with min key and swap(A[i], A[j])
    Min-Heapify(A, j)
```


This type of data structure naturally leads to a sorting algorithm. We can run a heapsort with the following pseudocode:

```
for i = n - 1 to 0:
    min-heapify(A, i) //this procedure takes  $O(n \log n)$ 
```

Our first legitimate run through heapsort brings us to this min-heap subtree, where  $30_2$  was swapped with its smaller child node 10:



After min-heapify iterates fully through the for loop, 10 "floats up" to the root and we obtain the following heap:



At every  $i$ , all subtrees rooted at  $j$  for  $j > i$  are valid and correct heaps. In a min-heap, a parent node is always less than or equal to its children nodes, and the root of the heap is always the minimum element. A parent is accessed by  $\frac{i}{2}$  and its children are given by  $2i + 1$  and  $2i + 2$ .

From here, we can use `extractMin()`, a function of a priority queue, which swaps the root node with the last node, and then returns the original root, and finally running `heapify` on the new root node. When this procedure is finished, we can call `extractMin()` again to fill up an array in sorted ascending order. The array that returns is such that:

10	20	30 <sub>2</sub>	30 <sub>1</sub>	40	50	70
----	----	-----------------	-----------------	----	----	----

With this example, we can also see that  $30_1$  and  $30_2$  are out of order and thus Heapsort is not a stable sorting algorithm.



**Question 2 (2 points):** Given an input array  $A$  of length  $n$  and a positive integer  $k > 0$ , design an algorithm that outputs the largest  $k$  elements in sorted order. Provide pseudocode and give a time-complexity analysis. You will get full credit if the time-complexity is  $O(n + k \log k)$ . (If you get anything less efficient, you only get 1 point.)

**Analysis:** Starting with the trivial solution, we could make a simple modification to Max-Heapsort, iterating our loop from  $n - 1$  to  $n - 1 - k$  instead of from  $n - 1$  to 0. Inside the loop, it would swap the root node, the maximum element, with the last node in the tree, then reducing the heap size by one and run Max-Heapify on the new root node. We could then return

the output in the form of  $A = \{A_{n-1}, \dots, A_{n-1-k}\}$  when the loop is finished. This algorithm would run in  $O(k \log n)$ . Here is some pseudocode:

```
for i = n - 1 to n - 1 - k:
    Heapsort(A) //this procedure takes  $O(n+k \log n)$ 
return A[n-1 ... n-1-k]
```

To achieve  $O(n + k \log k)$  time-complexity, I believe I would need to use two arrays  $A$  and  $B$ , and run max-heapify on both of them until they are max heaps. Then, let heap  $B$  act as a . This should run in  $O(n + k \log k)$ .



**Question 3 (1 point):** Prove that the number of keys stored in a 2-3 tree of height  $h$  is  $\Omega(2^h)$  and  $O(3^h)$ .

**Invariant:** The *invariant* of a 2-3 tree is that all of the children of every node have the same height  $h$ , and all leaf nodes lie on the same height  $h$ . Nodes also only ever have 0, 2, or 3 children – barring the nodes at level  $h$  which are all leaf nodes. Nodes also store one or two keys (also known as elements).

**Base case:** We can start with our base case, which is when  $h = 1$ . This is trivially true as level  $h = 0$  of the 2-3 tree is the root node, and at level  $h = 1$  there can be only two children nodes in the best-case scenario, and three children nodes in the worst-case scenario, so thus our base case  $\in \Omega(2^h) = \Omega(2^1)$  and  $\in O(3^h) = O(3^1)$ .

**Induction:** Our invariant holds true for  $h = 1, \dots, k$ . The keys, or elements, of a 2-3 tree are stored in each node, so our algorithm  $h \in \Omega(2^k)$  in our best-case scenario and  $\in O(3^k)$  in our worst-case scenario 2-3 tree. Now we can use induction to prove that it is true for  $h = 1, \dots, k, k + 1$ .

Nodes can only ever have two or three children, and nodes store the either one or two keys. This implies  $2^h$  nodes in the best-case with a runtime  $\in \Omega(2^h)$ , and  $3^h$  nodes in the worst-case with a runtime  $\in O(3^h)$ . If this is the case, then when  $h + 1 = k + 1$ , we have  $2^{h+1}$  nodes in the best-case with a runtime  $\in \Omega(2^{h+1})$  and  $3^{h+1}$  nodes in the worst-case with a runtime


$\in O(3^{h+1})$ .

Thus, the number of keys stored in a 2-3 tree with a height of  $h$  is  $\Omega(2^h)$  and  $O(3^h)$ .



**Question 4 (2 points):** Suppose you are given two 2-3 trees  $T_1$ ,  $T_2$ , and a value  $x$  such that all keys in  $T_1$  are less than  $x$ , and all keys in  $T_2$  are greater than  $x$ . Give an algorithm that constructs a single new  $T_0$  that has the union of keys in  $T_1$ ,  $T_2$ , and  $x$ . Give a running time analysis.

**Analysis:** First, we can insert()  $x$  into one of the smaller tree's leaves, which I will use as  $T_1$  in my example. First, we will need to search through the tree to find an appropriate leaf to insert  $x$  into. If a leaf already has one existing key, we simply just insert our new key  $x$  into the leaf. If the leaf has two existing keys, we still insert  $x$ . Finally, we run fix-overfull() on the node. This will take  $O(\log n)$  time. Here is what the result will look like:




Now that we have  $T_1 \cup x$  in a single 2-3 tree, we need to find a way to "merge" the unioned 2-3 tree with  $T_2$ . This union of  $T_1$  and  $x$  can also be treated as a single node, like this example:



Since a 2-3 tree is also technically a node, we can use this fact to insert()  $T_1 \cup x$  into one of  $T_2$ 's leaves and call fix-overfull() after, as fix-overfull()

accepts nodes as a parameter. This will take  $O(\log n)$  worst-case runtime. Here is a visualization of that merge if the keys in  $T_1 \cup x$  are less than the keys in  $T_2$  (it would go to the far right leaf if the keys in  $T_1 \cup x$  are greater than the keys in  $T_2$ ):



The 2-3 tree would be unbalanced with this addition, so we call `fix-overfull()` on the root node of  $T_1 \cup x$ , which would balance the 2-3 tree to satisfy the properties all three properties of a 2-3 tree.



**Question 5 (3 points):** Consider a binary search tree where keys are positive integers. Augment the tree to answer Range queries of the form: how many elements have key in the range  $[a, b]$ ? Thus, such a query is called by the function `Range( $a, b$ )`.

Provide pseudocode for Insert, Delete, and Range queries. Provide a running time analysis for all these queries, in terms on  $n$  (the number of nodes in the tree) and  $D$  (the maximum depth). (Hint: you might want to maintain subtree sizes at the nodes.)

**Analysis:** To solve this, I would have to envision a binary search tree that, for each node, would hold the number of nodes in a parent's subtree. Now we can start to perform our range query. It would recursively call itself, and if the node has a null key, return. If the node has a key that is less than or equal to  $a$ , everything to the left of it must be out of the range, so we add it to a counting variable. If the node has a key that is greater than or equal to  $b$ , everything to the right must be out of the range, so we add it to a counting variable. At the end of the recursion, we return the main root

of the entire tree's left and right subtree size to get a total tree size, and then subtract it from our counting variable. Here is some pseudocode (the runtime would be  $O(D+n)$ , and  $D \leq n$  so  $O(n)$ ):

```
count = 0;

Range(node, a,b): //this runs in  $O(n)$ 
    if (node.key == null) {return}
    if (node.key <= a){ count += node.leftsize, return }
    Range(node.left,a,b)
    if (node.key >= b){ count += node.rightsize, return }
    Range(node.right,a,b)
    return mainroot.lefttreesize+mainroot.righttreesize-count
```

We would need to update tree sizes as we traverse the tree in this environment. Here is some pseudocode to display how we would modify insert to handle keeping track of a subtree's size:

```
insert(node, val):
    if (node.key == null) { node.key = val }
    else if (val < node.key) {
        node.leftsize++
        insert(node.left, val)
    } else {
        node.rightsize++
        insert(node.right, val)
    }
```

Additionally, we would need a delete function that checks if there are no children, one child, or two children in the tree before deletion. This runtime of this would be  $O(n)$ . Here is some pseudocode to display how we would modify delete to handle keeping track of a subtree's size:

```
delete(node, val): //this is an  $O(n)$  operation
    if (node.key == val) {
        if (node.left.key == null && node.right.key == null) {
            node.key = null
        }
        if (node.left.key == null && node.right.key != null) {
```

```

        node.key = node.right.key
        node.right.key = null
    }
    if (node.left.key != null && node.right.key == null) {
        node.key = node.left.key
        node.left.key = null
    }
    if (node.left.key != null && node.right.key != null) {
        if (node.left.key > node.right.key) {
            node.key = node.left.key
            node.left.key = null
        } else {
            node.key = node.right.key
            node.right.key = null
        }
    }
} else if (val < node.key) {
    node.leftsize--
    delete(node.left, val)
} else {
    node.rightsize--
    delete(node.right, val)
}

```

■